
PHYSIGYM: Bridging Agent-Based Biological Simulation and Reinforcement Learning A State-Space Study on Tumor Microenvironment Control

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 This paper presents PHYSIGYM, a framework that integrates agent-based biological
2 simulation within standardised reinforcement learning environments. By connect-
3 ing the agent-based modelling framework PHYSICELL with the Gymnasium API,
4 we provide a flexible tool for exploring RL strategies that control *in silico* biolog-
5 ical processes. We demonstrate PHYSIGYM’s potential through a case study on
6 a simplified tumor microenvironment (TME) toy model, in which a Soft Actor-
7 Critic agent learns targeted drug-delivery policies that steer the microenvironment
8 toward an anti-tumoral state and, in the best episodes, achieve tumor elimination.
9 We further use this case study to systematically compare a family of observation
10 spaces from a partially observed macrophage-count baseline, through spatial scalar
11 summaries, to full multi-channel spatial images. Training on a structured rectangle
12 layout and testing on a held-out network-field layout, we find that spatially ex-
13 plicit image observations generalize markedly better than scalar summaries: image
14 modes reach a mean held-out test return of $\approx +30$ to $+32$, whereas spatial-scalar
15 summaries train comparably ($+43$) but collapse to ≈ 0 out of distribution, and the
16 macrophage-only POMDP baseline fails to learn at all (-32). This demonstrates
17 that PHYSIGYM’s flexible state-space interface directly impacts policy quality and
18 generalization in spatially structured environments. PHYSIGYM is open-source at
19 <https://anonymous.4open.science/r/PhysiGym-76A7>.

20 1 Introduction

21 Biological systems are complex, dynamic, and often difficult to study in real life. In silico modeling
22 aims to mimic these systems, providing a controlled environment for exploring specific aspects of
23 biological processes and testing hypotheses. In particular, agent-based models (ABMs) provide
24 a powerful framework for simulating biological systems where agents (e.g. cells) interact dynam-
25 ically based on predefined rules. These models are used in oncology and immunology to study
26 emerging behaviors in complex biological systems, with the potential to propose novel therapeutic
27 strategies [Johnson et al., 2025, Laubenbacher et al., 2024, Rocha et al., 2024].

28 Ordinary differential equations (ODEs) and agent-based models (ABMs) are widely used to simulate
29 the tumor microenvironment (TME). ODEs provide compact system-level descriptions, while ABMs
30 capture spatial and cellular heterogeneity at higher computational cost. For dynamic treatment
31 regimes (DTRs) [Gray and Pallavicini, 1982], however, simulation alone is insufficient: designing
32 adaptive, sequential interventions requires a control framework. Reinforcement learning (RL), where
33 an RL agent learns policies through interaction with an environment, offers such a framework [Barto
34 et al., 1989]. By modeling treatment regimes as a sequence of decisions, RL enables adaptive control

of biological simulations [Chakraborty and Murphy, 2014]. For efficient experimentation and policy learning, RL further requires a standardized interface. To address this, we introduce PHYSIGYM, an interface that connects PHYSICELL [Ghaffarizadeh et al., 2018], an ABM framework for multicellular simulations of tissues or cell cultures (tumor microenvironment, other disease tissues, or bacteria colonies), with Gymnasium [Towers et al., 2024], a widely adopted RL application programming interface (API). PHYSICELL enables flexible multiscale modeling of biological systems, coupled with a sophisticated underlying physics simulator (the BioFVM diffusion transport solver [Ghaffarizadeh et al., 2016]), while Gymnasium provides a structured framework for training and evaluating RL policies.

The first implementation of RL on an ABM framework based on BioFVM was proposed by Zade et al. [2020], which applied Q-learning [Watkins and Dayan, 1992] to optimize Temozolomide treatment regimes for patients with glioblastoma multiforme, a highly specific problem. More recently, Aif et al. [2025] used PHYSICELL as a complex (high-fidelity) simulation environment for tumor growth, combined with a simple (low-fidelity) RL training environment of coupled ordinary differential equations (ODE) to learn treatment strategies for solid tumors.

PHYSIGYM bridges PHYSICELL and Gymnasium to provide a reproducible and scalable framework for integrating agent-based simulations of biological systems with RL strategies. In the following, we present PHYSIGYM’s motivation and design, demonstrating how it directly connects ABMs and RL to control complex biological systems. In particular, we focus on the tumor microenvironment, a biological system made up of cancer cells and surrounding immune and stromal cells that engage in complex interactions and cross-talks, which is proven to be crucial for an effective response to immunotherapy, an important therapeutic approach in an increasing number of types of cancer.

To illustrate the framework, we build a simplified TME toy model and use it as an end-to-end RL case study. This model also serves as a controlled testbed for a practical question that arises in any spatially explicit biological environment: among the many observation spaces PHYSIGYM can expose (scalars, images, graphs), which ones help an RL agent learn an effective dynamic treatment regime? We find that image-based observations that preserve the spatial layout of cells and substrates outperform compact scalar summaries, both in cumulative return and in generalization to held-out spatial configurations. This is an application-level finding on a deliberately simplified model; the primary contribution of this work is PHYSIGYM itself.

Contributions. (1) PHYSIGYM: an open-source framework that bridges the PHYSICELL ABM and the Gymnasium RL API, turning any PHYSICELL model into an RL environment with minimal code changes. (2) A simplified TME toy model demonstrating an end-to-end RL case study in which an agent learns a spatial drug-delivery policy that drives the microenvironment toward an anti-tumoral state. (3) A systematic comparison of observation spaces on the toy model, showing that spatially explicit (image) representations help RL learn and generalize better than scalar summaries.

2 Methods

2.1 PHYSIGYM design and implementation

PHYSICELL is an open-source ABM framework written in C++ that models multicellular systems using classical mechanics and the BioFVM diffusive transport solver [Ghaffarizadeh et al., 2018, 2016]. Cells are agents governed by type-specific rule sets; substrates such as oxygen or cytokines diffuse through the domain. Intracellular signalling networks [Letort et al., 2018, Ponce-de Leon et al., 2023, Ruscone et al., 2024] and extracellular matrix fibres [Noël et al., 2024] can be added, yielding ABMs that are spatially explicit in 2D or 3D, off-lattice, and multiscale [Metzcar et al., 2019].

RL frames sequential decision-making as a Markov Decision Process (MDP) [Bellman, 1957, Puterman, 2014] with state space S , action space A , transition model T , and reward R . The objective is

$$\arg \max_{\pi \in \Pi} \mathbb{E} \left[\sum_{t=0}^{\tau} \gamma^t r_t \mid s_0, \pi \right].$$

Gymnasium [Towers et al., 2024] is the standard Python interface for MDPs, promoting modularity and code reuse across algorithms and environments.




Figure 1: PHYSIGYM dataflow. The PHYSICELL C++ loop is exposed to Python as three callables (`physicell_start/step/stop`); only `physicell_step` is modified per model (here, to inject `drug_1`). Three read functions (`get_cell`, `get_microenv`, `get_graph`) surface cell positions, substrate fields, and neighbourhood graphs. A user subclasses `CorePhysiCellEnv` to assemble these into a Gymnasium observation/reward, and any RL algorithm drives the environment through the standard `step/reset` API.

PHYSIGYM connects these two ecosystems by extending the Python interpreter [van Rossum, 1995, Python Software Foundation, 2024] with PHYSICELL. The C++ simulation loop is split into three Python-callable functions: `physicell_start`, `physicell_step`, and `physicell_stop`. For most models only `physicell_step` needs to be modified (e.g. to inject a drug). Three read functions `get_cell`, `get_microenv`, and `get_graph` return cell positions, substrate concentration fields, and cell neighbourhood graphs respectively; all are implemented in `physicellmodule.cpp`.

The interface between PHYSICELL and Gymnasium uses PHYSICELL’s built-in *user parameter* and *custom data* constructs, configurable through PhysiCell Studio [Heiland et al., 2024] without recompilation. Users subclass `CorePhysiCellEnv` in `physicell_model.py` to specify observation space, action space, and reward; the base class handles all model-independent Gymnasium bookkeeping. Any existing PHYSICELL model can therefore become a PHYSIGYM environment with minimal code changes. PHYSIGYM is tested on Linux, WSL, and macOS, with continuous integration via GitHub Actions. Documentation and tutorials are at <https://anonymous.4open.science/r/PhysiGym-76A7/man>.

Figure 1 shows the resulting dataflow, and Listing 2 gives a complete minimal environment: a user overrides only the observation, action, and reward, while the base class handles the PHYSICELL `start/step/stop` loop and all Gymnasium bookkeeping.

2.2 Implementation of a Tumor Microenvironment Toy Model

We developed a simplified tumor microenvironment (TME) *toy model* to illustrate the potential uses of PHYSIGYM. The model is deliberately small and computationally cheap so that we can train RL agents to convergence many times over (across observation spaces and seeds), while still retaining the spatial, multicellular, and chemically coupled structure that distinguishes agent-based models (ABMs) from ODE surrogates. It includes three cell types, five diffusible substrates, and a set of simple agent-based rules. This toy model highlights how the RL framework can control tumor progression by modulating the whole microenvironment, rather than solely targeting cancer-cell killing.

```

from physigym.envs import CorePhysiCellEnv
from gymnasium import spaces
import numpy as np

class ModelPhysiCellEnv(CorePhysiCellEnv):
    # Action: localized drug injection (dose, x, y, radius), all normalized.
    def get_action_space(self):
        box = lambda: spaces.Box(0.0, 1.0, (1,), np.float32)
        return spaces.Dict({"drug_1_dose": box(), "drug_1_x": box(),
                           "drug_1_y": box(), "drug_1_radius": box()})

    # Observation: e.g. per-cell-type 64x64 image channels (img_mc_cells).
    def get_observation_space(self):
        return spaces.Box(0, 255, (self.n_channels, 64, 64), np.uint8)

    def get_observation(self):
        # get_cell / get_microenv surface ABH state from the C++ core.
        return self.render_cell_channels() # -> (C, 64, 64) uint8

    # Reward: tumor suppression minus a drug-cost penalty (Eq. in paper).
    def get_reward(self):
        dc = (self.count_prev - self.count) / self.growth_norm
        return self.w_cell * dc - self.w_dose * self.get_dose_spent()

    def get_terminated(self):
        return self.count == 0 # tumor eradicated

```

Figure 2: A complete PHYSIGYM environment: subclass `CorePhysiCellEnv` and override the observation, action, and reward. The base class handles all model-independent Gymnasium book-keeping (reset/step, the PHYSICELL start/step/stop loop, and the `get_cell/get_microenv` read-back), so no C++ recompilation is needed to expose a new state space. (Condensed from `physicell_model.py`.)

110 **Domain and Dynamics.** The continuous-space environment is simulated on a $64 \times 64 \mu\text{m}$ 2D
111 domain. The RL gym step interval is $\Delta t = 15$ min, with episodes running up to 10,080 min (7
112 simulated days).

113 **Entity Dynamics & The Control Challenge.** The environment formulates a highly non-linear,
114 indirect control problem. The agent’s action (deploying a localized therapeutic, `drug_1`) has zero
115 direct cytotoxicity. Instead, the agent must learn to manipulate the environment’s underlying transition
116 dynamics by reversing a spatial immunosuppression loop. The system comprises three interacting
117 cellular populations whose movements are governed by biochemical spatial gradients (chemotaxis):

- 118 • **Tumor Cells (The Target):** Immobile entities that continuously divide. They actively
119 modify the environment by secreting a chemical signal (`tumor_molecule`) that suppresses
120 the local immune system.
- 121 • **Macrophages (The Mediators):** Mobile, immortal entities that navigate toward the tumor.
122 They exhibit state-dependent behavior: the tumor’s chemical signal successfully “hijacks”
123 them, transforming them from an anti-tumoral to a pro-tumoral state.
- 124 • **T-cells (The Effectors):** Mobile, immortal entities capable of triggering tumor apopto-
125 sis. However, their spatial navigation is strongly repelled by the presence of pro-tumoral
126 macrophages.

127 This setup creates a hostile, self-reinforcing feedback loop: the tumor protects itself by converting
128 macrophages, which in turn dynamically block T-cells from reaching the target. As illustrated in
129 Figure 3, the RL agent must maximize reward by learning to precisely deploy its drug to “re-program”
130 the macrophages back to an anti-tumoral state. This spatial intervention remodels the chemical
131 gradients, clearing a path for T-cells to navigate to the tumor and execute their killing function.

132 2.3 Implementation of an RL approach to control the TME

133 Our overall objective is to achieve complete tumor eradication while minimizing drug exposure. We
134 describe the action space A , the reward R , and the training and test distributions. The observation




Figure 3: Tumor Microenvironment (TME) signalling network. Green solid arrows denote production, activation, or attraction, while red solid arrows indicate suppression or repulsion. The dashed pathways represent the intervention of the agent-controlled drug (drug_1), highlighting its role as an immune-modulator that repolarizes macrophages to indirectly drive T-cell mediated tumor apoptosis.

space S is the component we vary across experiments: in Section 3 we compare a family of observation modes spanning a partially observed scalar baseline, spatial scalar summaries, and multi-channel spatial images to assess whether spatially explicit representations help the agent learn and generalize on this toy model.

Action space. At each gym step the agent selects $a_t = (\text{dose}, x, y, r) \in [0, 1]^4$, where dose is the drug concentration applied, (x, y) is the injection centre (normalised domain coordinates), and r is the injection radius (normalised by the domain diagonal).

Reward. The reward is a normalised tumor-suppression signal minus a weighted drug-cost penalty:

$$r_t = w_{\text{cell}} \frac{c_{t-1} - c_t}{c_{t-1}(e^{\lambda \Delta t} - 1)} - w_{\text{dose}} d_t, \quad (1)$$

where c_t is the alive tumor-cell count at step t , λ is the intrinsic growth rate (from `PhysiCell_settings.xml`), d_t is the normalised dose spent, $w_{\text{cell}} = 0.3$, and $w_{\text{dose}} = 2$. Raising the dose weight from 1 to 2 relative to earlier experiments strengthens the pressure to spend drug sparingly, so the agent is rewarded for converting each unit of dose into effective macrophage repolarisation rather than for blanket dosing. The d_t term penalises drug use directly, keeping the learning signal clean without conflating temporal dosing strategy with total exposure.

Training and Test Distributions

To ensure the model learns robust policies rather than overfitting to specific spatial configurations, we train and test on two *structurally distinct* spatial distributions, so that test performance measures out-of-distribution (OOD) generalisation rather than memorisation of the training geometry:

- **Rectangle (training):** Each cell type is placed uniformly in a narrow vertical strip, with non-overlapping strip positions (Appendix A). This is a highly structured, low-entropy layout.
- **Network-field (held-out test):** Initial positions are sampled from a spatially correlated random field (Appendix A), producing irregular, clustered configurations never seen during training.

Training on the simple, regular rectangle layout and testing on the richer, irregular network-field is a demanding transfer: a policy that merely exploits the fixed strip geometry of the rectangle cannot solve the network-field, so the test return isolates whether an observation space encodes *transferable* structure. Figure 8 illustrates representative examples of both distributions.

163 2.4 State space S : observation spaces compared

164 In a tumor microenvironment there are different cell types and different molecules/chemical com-
 165 pounds (substrates). The simulator can expose several kinds of data: some carry explicit spatial
 166 information, others are spaceless summaries (e.g. a small vector of global statistics). Because PHYSI-
 167 GYM can build all of these from the same underlying simulation, we use the toy model to study how
 168 the choice of observation space affects the dynamic treatment regime that RL learns, comparing nine
 169 modes on cumulative return and on generalization to unseen spatial configurations.

170 We compare observation modes spanning a spectrum from a partially observed scalar baseline to
 171 full multi-channel spatial images. Let $N_c = 3$ (tumor, T-cell, macrophage) and $N_s = 5$ substrates.
 172 Several modes optionally split the macrophage population into its $M1$ (anti-tumoral) and $M2$ (pro-
 173 tumoral) polarisation states, which we denote with the suffix `m1m2`; the classification rule is given
 174 below.

175 **Macrophage polarisation (M1/M2).** Macrophages do not carry an explicit polarisation label in
 176 the simulator; their state is implicit in the local chemical environment. We recover it from the
 177 microenvironment: for each alive macrophage we look up the two competing substrate fields at its
 178 nearest voxel and classify it as $M2$ (pro-tumoral) if $\text{conc}_{\text{pro_tumoral}} > \text{conc}_{\text{anti_tumoral}}$ at that voxel,
 179 and $M1$ (anti-tumoral) otherwise. This makes the pro-/anti-tumoral balance the very quantity the
 180 agent must flip observable as a count (scalars modes) or as two spatial channels (`img` modes),
 181 rather than leaving it latent in the substrate fields.

182 Scalar Observation Spaces

183 **scalars_macrophages (POMDP baseline):** Alive count per cell type, with the macrophage slot
 184 split into M1 and M2 counts; each is normalised to the initial population ($s = c/c_{\text{init}} - 1$). This mode
 185 exposes only *how many* macrophages are polarised each way and carries no spatial information at
 186 all neither cell positions nor substrate maps. It is therefore a strongly partially observed (POMDP)
 187 baseline: the true state that governs the transition dynamics (where cells and gradients are) is hidden,
 188 and the agent must act on an aggressively compressed summary.
 189 *Space:* $\mathbb{R}^{N_c+1}$

190 **spatial_scalars_cells_substrates:** Per-type normalised cell counts and per-substrate max-
 191 ima, augmented with 6 k -means spatial statistics per cell type (cluster occupancy and centroids),
 192 giving a compact spatial summary without full images.
 193 *Space:* $\mathbb{R}^{N_c+N_s+6kN_c}$

194 **spatial_scalars_cells_m1m2 / ..._substrates_m1m2:** As above but with the macrophage
 195 population split into M1/M2 for both the scalar counts and the k -means spatial features
 196 (..._substrates_m1m2 additionally keeps the substrate maxima).
 197 *Space:* $\mathbb{R}^{(N_c+1)+[N_s]+6k(N_c+1)}$

198 **spatial_scalars_cells_spatial_no_scalars_substrates_m1m2:** M1/M2 cell counts
 199 plus k -means spatial features for *both* cell types and substrate fields (spatially resolved substrate
 200 summaries in place of scalar maxima).
 201 *Space:* $\mathbb{R}^{(N_c+1)+6kN_s+6k(N_c+1)}$

202 Image-Based Observation Spaces

203 **img_mc_cells (I1):** One 64×64 channel per cell type.
 204 *Space:* $\mathbb{R}^{N_c \times 64 \times 64}$ (uint8)

205 **img_mc_cells_m1m2:** I1 with the macrophage channel split into separate M1 and M2 spatial
 206 maps.
 207 *Space:* $\mathbb{R}^{(N_c+1) \times 64 \times 64}$ (uint8)

208 **img_mc_cells_substrates (I2):** I1 concatenated with all five substrate concentration maps,
 209 providing direct spatial feedback on drug distribution and the polarising factors.
 210 *Space:* $\mathbb{R}^{(N_c+N_s) \times 64 \times 64}$ (uint8)

211 `img_mc_cells_substrates_m1m2` (**full information**): I2 with the additional M1/M2
212 macrophage channels. This mode carries every signal the others expose cell positions, substrate fields,
213 and the explicit polarisation split and serves as the near-fully-observed upper reference against which
214 the compressed modes are compared.
215 *Space*: $\mathbb{R}^{(N_c+N_s+2)\times 64\times 64}$ (uint8)

216 An observability spectrum

217 Together these modes trace an observability spectrum, from `scalars_macrophages` a POMDP that
218 sees only polarisation counts to `img_mc_cells_substrates_m1m2`, which spatially resolves every
219 state variable relevant to the control problem. The indirect drug mechanism (Section 2.2) requires
220 multi-step credit assignment: the agent injects drug, macrophages repolarise over several steps,
221 anti-tumoral gradients shift, T-cells migrate, and only then do tumor cells die. Scalar summaries
222 cannot tell the agent whether an injection reached macrophages, whether polarisation occurred, or
223 where T-cells will concentrate next; only the per-pixel substrate and M1/M2 channels of the image
224 modes make the polarisation state and drug distribution spatially legible at each step.

225 3 Results

226 3.1 Aggregated performance

227 Table 1 summarises final performance (seeds averaged). Figures 4 and 6 show learning curves vs.
228 cumulative environment steps with mean $\pm$ std shading across seeds.

Table 1: Performance at end of training (train = rectangle, test = network-field, seeds averaged). $\text{Train}_\mu/\text{Test}_\mu$: mean EWMA-50 return. $\text{Train}_\sigma/\text{Test}_\sigma$: std across seeds; n = number of seeds. This is the curated main set; the remaining m1m2 scalar variants appear in Appendix C, and 95% bootstrap confidence intervals on the mean in Appendix D. RAND is a random-action policy (5 seeds), included as a learning-free reference: modes below it fail to beat random behaviour.

ID	Mode	n	Train_μ	Train_σ	Test_μ	Test_σ
I2m	<code>img_mc_cells_substrates_m1m2</code>	5	+40.8	8.5	+32.5	9.4
I1	<code>img_mc_cells</code>	6	+30.5	8.7	+32.4	10.5
I2	<code>img_mc_cells_substrates</code>	7	+42.5	6.6	+30.4	7.4
S3s	<code>spatial_scalars_cells_substrates</code>	5	+42.9	18.8	−0.0	9.1
RAND	random policy (baseline)	5	−18.4	4.0	−25.3	3.0
POMDP	<code>scalars_macrophages</code>	5	−20.4	8.3	−31.8	2.3

229 3.2 Key findings

230 **Generalization, not training fit, separates the modes.** On the training rectangle layout the spatial-
231 scalar summary S3s is fully competitive its training return (+42.9) matches the best image mode I2
232 (+42.5). The distinction appears only out of distribution: on the held-out network-field, the three
233 image modes retain $\approx +30$ to +32 test return, while S3s collapses to −0.0. Training fit alone would
234 have ranked the scalar summary among the best; only the OOD test reveals that it has overfit the
235 fixed strip geometry of the rectangle. This separation is statistically supported: the image modes’
236 95% bootstrap confidence intervals on the mean test return lie entirely above +24 and do not overlap
237 those of any spatial-scalar mode, all of which contain zero (Appendix D).

238 **Spatially explicit observations transfer; scalar summaries do not.** All three image modes
239 `img_mc_cells` (I1, +32.4), `img_mc_cells_substrates` (I2, +30.4), and the full-information
240 `img_mc_cells_substrates_m1m2` (I2m, +32.5) generalize to the network-field, and the gaps
241 among them are within one seed-standard-deviation. That I1 already generalizes well indicates that,
242 in this toy model, preserving the *spatial layout* of cells is the primary driver of transfer; the extra
243 substrate and M1/M2 channels of I2/I2m do not further separate on average, though I2m gives the
244 single best test return.




Figure 4: Training return on the rectangle layout (EWMA-50, mean $\pm 1\sigma$ across seeds) vs. cumulative *simulator* steps ($\times 10^5$; each of the 10^5 agent decisions is held for 6 simulator steps, giving a 6×10^5 budget). All modes learn on the training layout except the scalars_macrophages POMDP baseline, which flatlines near -20 . Spatial-scalar summaries (S3s) train as fast as, or faster than, the image modes here the separation only appears out of distribution (Figure 6).

The POMDP baseline is no better than acting randomly. scalars_macrophages, which observes only M1/M2 counts and no spatial information, fails to improve on either distribution (train -20.4 , test -31.8) and, crucially, does *not* beat a random-action policy (train -18.4 , test -25.3 ; RAND in Table 1). Its test return is if anything below random, and the two bootstrap CIs barely overlap (Appendix D): the macrophage-count signal is not merely uninformative but actively misleading, steering dosing worse than chance. Knowing *how many* macrophages are polarised without knowing *where* the drug, cells, and gradients are is insufficient to solve the indirect, spatially mediated control problem a concrete illustration of how observation-space choice can render an otherwise learnable task effectively unlearnable. By contrast, every spatial-scalar mode, though it fails to generalize, still clears the random baseline on test (≥ 0 vs. -25.3), confirming they learn *some* transferable structure just far less than the image modes.

On seed robustness. Unlike a cleaner earlier run, we do not observe a systematic seed-variance advantage for any family on this benchmark: end-of-training test std is 7–12 for the image modes and 9 for S3s (Table 1, Figure 5). The POMDP baseline’s low std (2.3) reflects consistent failure rather than robustness.

3.3 The role of the M1/M2 split and spatial scalars

The full mode family (Appendix C, Table 3) lets us isolate the effect of the explicit M1/M2 polarisation split. Adding it to the cell image (I1 $\rightarrow$ img_mc_cells_m1m2) does not improve average OOD test return ($+32.4 \rightarrow +26.0$), and adding it on top of the full image (I2 $\rightarrow$ I2m) leaves it essentially unchanged ($+30.4 \rightarrow +32.5$): once the substrate fields are spatially available, the CNN can already read out polarisation from the pro_tumoral/anti_tumoral channels, so the pre-computed split is redundant. Among the spatial-scalar modes, none generalize (all within $[-0.0, +9.5]$ test return) regardless of whether they carry the M1/M2 split or spatially resolved substrate summaries the common failure is the loss of per-cell spatial layout, which the k -means summaries do not preserve.




Figure 5: Standard deviation of return across seeds vs. simulator steps: training on rectangle (left) and test on network-field (right). Inter-seed variability is broadly comparable across the learning modes and fluctuates over training, so we do not claim a systematic variance advantage for any single family; the POMDP baseline has the lowest test variance simply because it converges to a uniformly poor policy. At end of training the image modes reach test return $\approx +30$ with std $\approx 7\text{--}12$ (Table 1).

Appendix E shows a matched example episode comparing an image mode, a spatial-scalar mode, and the POMDP baseline rolled out from an identical network-field initial condition, illustrating the failure mode that summarises the comparison: without a spatially explicit representation the agent cannot target the drug precisely enough to convert dose into effective macrophage repolarisation on an unseen layout the image mode reaches near-eradication (3 tumor cells) while the scalar and POMDP modes leave 70 and 130 cells respectively.

4 Discussion

The central outcome of this work is PHYSIGYM itself: a framework that lets RL agents control a full agent-based TME through a standardised interface, without recompilation, and with arbitrary observation and action spaces. The toy-model case study shows the framework end-to-end an agent learns a spatial drug-delivery policy that steers the microenvironment toward an anti-tumoral state and, in the best episodes, eradicates the tumor.

The toy model also produces a clear result on observation spaces, and crucially one that only surfaces under distribution shift. On the training rectangle layout, spatial-scalar summaries are fully competitive with image observations; it is on the held-out network-field that spatially explicit (image) observations pull ahead, generalizing to $\approx +30$ test return while the scalar summaries collapse to ≈ 0 . The CNN in the image modes can compute cross-channel co-localisation statistics (e.g. whether drug₁ and macrophages overlap at the same location) directly from the pixel layout statistics that transfer across geometries whereas scalar modes that encode k -means cluster centroids and absolute positions overfit the fixed strip geometry of the training layout. Notably, once cell layout is spatially available, preserving the substrate fields (I2) or pre-computing the M1/M2 split (I2m) does not further improve average transfer over the cell-only image (I1) on this toy model, though the full-information mode gives the single best test return.

The overfit is a drug-aiming failure, measurable in the actions. The claim above is not merely architectural: it is visible directly in the policies’ behaviour. For every logged episode we compute the mean distance between the drug-injection centre (action_x, action_y) and the tumour centroid (from the episode initial conditions), averaged over the steps on which a dose is actually applied, and compare this aiming error between the training rectangle and the held-out network-field (Fig. 7). On the training layout all six modes aim equally well (mean normalised error $\approx 0.34\text{--}0.36$ for both families). Under the geometry shift, image modes hold their aim or slightly improve it (I2: $0.359 \rightarrow 0.343$; I2m: $0.352 \rightarrow 0.333$), whereas every scalar mode’s aim *degrades* (S3s: $0.341 \rightarrow 0.414$; S3m:




Figure 6: Held-out test return on the network-field layout (mean $\pm 1\sigma$ across seeds, EWMA-50) vs. cumulative simulator steps. The image modes (I2m, I2, I1) climb steadily to $\approx +30$, whereas the spatial-scalar summary (S3s) which trained *best* (Figure 4) collapses to ≈ 0 , revealing that it overfits the fixed rectangle geometry rather than learning transferable structure. The dashed line marks the random-policy baseline (test ≈ -25 , shaded $\pm 1\sigma$); the `scalars_macrophages` POMDP baseline tracks at or below it throughout, i.e. never beats random. Generalization, not training fit, is where spatially explicit observations separate from scalar summaries.

0.348 $\rightarrow$ 0.406) and its step-to-step variance roughly doubles. In other words, the scalar policies learn *where* the rectangle tumour tends to sit rather than *how to find* the tumour, so when the network-field moves it they inject into empty tissue; the median final tumour count on test rises accordingly (image ≈ 73 –99 vs. scalar 107–136). This is the concrete mechanism behind the aggregate return gap: the generalisation failure is an aiming failure induced by position-absolute features.

We emphasise that this is an application-level observation on a deliberately simplified model, not a universal claim about biological control. In more complex or realistic models or with different reward designs, action spaces, or RL algorithms the relative value of cell versus substrate, and scalar versus image, may differ. PHYSIGYM is precisely the tool that makes such questions answerable in a reproducible way.

Limitations. The toy model is intentionally small ($64 \mu\text{m}$, 2D) and cheap to simulate. Scaling to 3D or larger domains will increase the cost of image observations. All experiments use a single RL algorithm (SAC); alternative algorithms or model-based methods may interact differently with observation spaces. Each mode is trained with a modest number of seeds ($n = 5$ –7), unevenly distributed across modes; we mitigate this with bootstrap confidence intervals (Appendix D) rather than relying on raw seed-standard-deviation, but the conclusions would be strengthened by more seeds and by testing transfer in both directions between layouts. We also study a single train/test layout pair (rectangle $\rightarrow$ network-field); the reverse direction is left to future work. Complex biological models and/or 3D models are left to future work.

5 Conclusion

We presented PHYSIGYM, an open-source framework that bridges the PHYSICELL ABM with Gymnasium RL, enabling direct policy learning in spatially explicit biological simulations. Using a TME model with an indirect drug mechanism, we showed that spatially explicit image observations




Figure 7: Drug-aiming alignment under geometry shift. Bars show the mean normalised distance between the drug-injection centre and the tumour centroid over dosed steps (lower is better aiming); error bars are the across-episode standard deviation. On the training rectangle (blue) image and scalar observations aim equally well. On the held-out network-field (red) image modes retain their aim while every scalar mode degrades with markedly higher variance, exposing the overfit as a failure to re-target the tumour rather than a loss of dosing ability.

are the critical state representation for *generalization*: training on a structured rectangle layout and testing on a held-out network-field, image modes retain $\approx +30$ to $+32$ test return while spatial-scalar summaries competitive during training collapse to ≈ 0 out of distribution, and a macrophage-count-only POMDP baseline fails to learn at all. The gap is a generalization phenomenon: what separates the representations is not training fit but whether the encoded structure transfers across spatial configurations. PHYSIGYM provides the infrastructure to study such questions in a principled, reproducible manner, and we anticipate it will serve as a testbed for both computational biology and RL communities.

References

- S. Aif, M. Eiche, N. Appold, E. Fischer, T. Citak, and J. Kayser. Reinforcement failing guides the discovery of emergent spatial dynamics in adaptive tumor therapy. Apr. 2025. doi: 10.1101/2025.04.08.647768.
- A. G. Barto, R. S. Sutton, and C. Watkins. *Learning and sequential decision making*, volume 89. University of Massachusetts Amherst, MA, 1989.
- R. Bellman. A markovian decision process. *Journal of mathematics and mechanics*, pages 679–684, 1957.
- B. Chakraborty and S. A. Murphy. Dynamic treatment regimes. *Annual review of statistics and its application*, 1(1):447–464, 2014.
- L. Espeholt, H. Soyer, R. Munos, K. Simonyan, V. Mnih, T. Ward, Y. Doron, V. Firoiu, T. Harley, I. Dunning, et al. Impala: Scalable distributed deep-rl with importance weighted actor-learner architectures. In *International conference on machine learning*, pages 1407–1416. PMLR, 2018.
- A. Ghaffarizadeh, S. H. Friedman, and P. Macklin. Biofvm: an efficient, parallelized diffusive transport solver for 3-d biological simulations. *Bioinformatics*, 32(8):1256–1258, 2016.
- A. Ghaffarizadeh, R. Heiland, S. H. Friedman, S. M. Mumenthaler, and P. Macklin. PhysiCell: An open source physics-based cell simulator for 3-d multicellular systems. *PLOS Computational Biology*, 14(2):e1005991, feb 2018. doi: 10.1371/journal.pcbi.1005991.

349 J. W. Gray and M. G. Pallavicini. Ara-c scheduling: Theoretical and experimental considerations.
350 *Medical and Pediatric Oncology*, 10(S1):93–108, 1982. ISSN 1096-911X. doi: 10.1002/mpo.
351 2950100711.

352 T. Haarnoja, A. Zhou, K. Hartikainen, G. Tucker, S. Ha, J. Tan, V. Kumar, H. Zhu, A. Gupta,
353 P. Abbeel, et al. Soft actor-critic algorithms and applications. *arXiv preprint arXiv:1812.05905*,
354 2018.

355 R. Heiland, D. Bergman, B. Lyons, G. Waldow, J. Cass, H. L. da Rocha, M. Ruscone, V. Noël, and
356 P. Macklin. Physicell studio: a graphical tool to make agent-based modeling more accessible.
357 *Gigabyte*, 2024:gigabyte128, 2024.

358 J. A. Johnson, D. R. Bergman, H. L. Rocha, D. L. Zhou, E. Cramer, I. C. Mclean, Y. W. Dance,
359 M. Booth, Z. Nicholas, T. Lopez-Vidal, A. Deshpande, R. Heiland, E. Bucher, F. Shojaeian,
360 M. Dunworth, A. Forjaz, M. Getz, I. Godet, F. Kurtoglu, M. Lyman, J. Metzcar, J. T. Mitchell,
361 A. Raddatz, J. Solorzano, A. Sundus, Y. Wang, D. G. DeNardo, A. J. Ewald, D. M. Gilkes, L. T.
362 Kagohara, A. L. Kiemen, E. D. Thompson, D. Wirtz, L. D. Wood, P.-H. Wu, N. Zaidi, L. Zheng,
363 J. W. Zimmerman, J. M. Phillip, E. M. Jaffee, J. W. Gray, L. M. Coussens, Y. H. Chang, L. M.
364 Heiser, G. L. Stein-O’Brien, E. J. Fertig, and P. Macklin. Human interpretable grammar encodes
365 multicellular systems biology models to democratize virtual cell laboratories. *Cell*, July 2025.
366 ISSN 0092-8674. doi: 10.1016/j.cell.2025.06.048.

367 R. Laubenbacher, F. Adler, G. An, F. Castiglione, S. Eubank, L. L. Fonseca, J. Glazier, T. Helikar,
368 M. Jett-Tilton, D. Kirschner, et al. Toward mechanistic medical digital twins: some use cases in
369 immunology. *Frontiers in Digital Health*, 6:1349595, 2024.

370 G. Letort, A. Montagud, G. Stoll, R. Heiland, E. Barillot, P. Macklin, A. Zinovyev, and L. Calzone.
371 PhysiBoSS: a multi-scale agent-based modelling framework integrating physical dimension and
372 cell signalling. *Bioinformatics*, 35(7):1188–1196, aug 2018. doi: 10.1093/bioinformatics/bty766.

373 X.-Y. Liu, D. Bodaghi, Q. Xue, X. Zheng, and J.-X. Wang. Asynchronous parallel reinforcement
374 learning for optimizing propulsive performance in fin ray control. *Engineering with Computers*, 41
375 (4):2201–2218, 2025.

376 J. Metzcar, Y. Wang, R. Heiland, and P. Macklin. A review of cell-based computational modeling in
377 cancer biology. *JCO Clinical Cancer Informatics*, pages 1–13, Dec. 2019. ISSN 2473-4276. doi:
378 10.1200/cci.18.00069.

379 V. Mnih, A. P. Badia, M. Mirza, A. Graves, T. Lillicrap, T. Harley, D. Silver, and K. Kavukcuoglu.
380 Asynchronous methods for deep reinforcement learning. In *International conference on machine*
381 *learning*, pages 1928–1937. PmLR, 2016.

382 V. Noël, M. Ruscone, R. Shuttleworth, and C. K. Macnamara. Physimess - a new physicell addon
383 for extracellular matrix modelling. *Gigabyte*, 2024, Oct. 2024. ISSN 2709-4715. doi: 10.46471/
384 gigabyte.136.

385 K. O’shea and R. Nash. An introduction to convolutional neural networks. *arXiv preprint*
386 *arXiv:1511.08458*, 2015.

387 M. Ponce-de Leon, A. Montagud, V. Noël, A. Meert, G. Pradas, E. Barillot, L. Calzone, and
388 A. Valencia. Physiboss 2.0: a sustainable integration of stochastic boolean and agent-based
389 modelling frameworks. *npj Systems Biology and Applications*, 9(1):54, 2023.

390 M. L. Puterman. *Markov decision processes: discrete stochastic dynamic programming*. John Wiley
391 & Sons, 2014.

392 Python Software Foundation. Extending and embedding the python interpreter. <https://docs.python.org/3/extending/embedding.html>, 2024. Accessed: 2024-02-07.

393
394 L. H. Rocha, B. Aguilar, M. Getz, I. Shmulevich, and P. Macklin. A multiscale model of immune
395 surveillance in micrometastases gives insights on cancer patient digital twins. *npj Systems Biology*
396 *and Applications*, 10(1), Dec. 2024. ISSN 2056-7189. doi: 10.1038/s41540-024-00472-z.

- 397 M. Ruscone, A. Checcoli, R. Heiland, E. Barillot, P. Macklin, L. Calzone, and V. Noël. Building
 398 multiscale models with physiboss, an agent-based modeling tool. *Briefings in Bioinformatics*, 25
 399 (6), Sept. 2024. ISSN 1477-4054. doi: 10.1093/bib/bbae509.
- 400 M. Towers, A. Kwiatkowski, J. Terry, J. U. Balis, G. De Cola, T. Deleu, M. Goulão, A. Kallinteris,
 401 M. Krimmel, A. KG, R. Perez-Vicente, A. Pierré, S. Schulhoff, J. J. Tai, H. Tan, and O. G. Younis.
 402 Gymnasium: A standard interface for reinforcement learning environments, 2024.
- 403 G. van Rossum. Extending and embedding the python interpreter. *CWI Report CS-R9527*, May 1995.
 404 URL <https://gvanrossum.github.io/Publications.html>.
- 405 C. J. Watkins and P. Dayan. Q-learning. *Machine learning*, 8:279–292, 1992.
- 406 A. E. Zade, S. S. Haghighi, and M. Soltani. Reinforcement learning for optimal scheduling of
 407 glioblastoma treatment with temozolomide. *Computer methods and programs in biomedicine*, 193:
 408 105443, 2020.

409 A Generation of Initial Conditions

410 **Rectangle (training layout).** Each cell type is placed uniformly in a vertical strip of width 10% of
 411 the x -axis, with non-overlapping strip positions sampled from $[0.1, 0.2]$, $[0.3, 0.5]$, $[0.6, 0.8]$ (shuffled
 412 across types). This is a highly regular, low-entropy layout.

413 **Network-field (held-out test layout).** White Gaussian noise $\eta(\mathbf{x}) \sim \mathcal{N}(0, 1)$ is convolved with a
 414 Gaussian kernel G_σ ($\sigma = \ell_c/\sqrt{2}$, $\ell_c = 45 \mu\text{m}$) to produce a correlated field F . An additional local
 415 noise component $\tilde{\xi}$ filtered at width $\ell_c/3$ is added with small weight α to introduce local heterogeneity.
 416 The field is normalised to $[0, 1]$ and cells are sampled from positions with $\hat{F} > 0.55$, with probability
 417 proportional to $\hat{F}$, producing irregular, clustered configurations that are structurally distinct from the
 418 training rectangle.




Figure 8: Representative examples of the training (rectangle) and held-out test (network-field) spatial layouts.

B Training algorithm and setup

To address our control objective maximizing cancer-cell eradication while minimizing drug administration we employ Soft Actor-Critic (SAC) [Haarnoja et al., 2018], a state-of-the-art RL algorithm for continuous control. SAC is well suited to our setting due to its high sample efficiency, and as a deep RL method it accommodates different observation types, such as image-based representations, through Convolutional Neural Networks (CNNs) [O’shea and Nash, 2015]. A key challenge is the computational cost of the simulator; to alleviate this bottleneck we adopt an asynchronous parallel training framework [Mnih et al., 2016, Liu et al., 2025]. A GPU-resident learner updates policy and value networks asynchronously while 13 CPU-resident environment processes generate experience in parallel.

Compute. Each run trains a single observation mode for one seed to 10^5 agent decisions (6×10^5 simulator steps) in approximately 5 hours of wall-clock time on a single workstation (one GPU-resident learner with 13 CPU environment workers). The asynchronous design decouples the GPU learner from the comparatively slow PHYSICELL simulator, so throughput is limited by the CPU environment workers rather than by gradient updates; this is what makes the full observation-space sweep (nine modes $\times$ 5–7 seeds) tractable on a single machine.

Action repeat. Each action selected by the agent is held fixed for 6 consecutive simulator steps (*action repeat* of 6) before the next decision. This serves two purposes. First, it stabilises learning: with a 15-minute decision interval on a slow, indirect biological process (drug injection $\rightarrow$ macrophage repolarisation $\rightarrow$ gradient shift $\rightarrow$ T-cell migration $\rightarrow$ tumor death), consecutive raw simulator steps are highly correlated and a fresh decision at every step injects high-frequency jitter into the drug field without changing the underlying dynamics; committing to an action for several steps lets its effect materialise before it is revised, which empirically yields smoother dose trajectories and lower-variance value targets. Second, it improves simulator throughput per decision: the 10^5 agent decisions per run correspond to 6×10^5 simulator steps of experience, so the effective interaction budget (and all step axes in Figures 4–6) is 6×10^5 . Action repeat is a standard device for temporally extended control and is well motivated here by the slow timescale of the controlled process relative to the simulator step. The specific value $AR = 6$ is selected by a sensitivity sweep over $AR \in \{1, \dots, 8\}$ (Appendix B.4, Figure 9): the action-smoothness cost drops steeply and plateaus by $AR \approx 6$, and the composite control objective reaches its plateau over $AR = 5$ –8, so $AR = 6$ is the smallest repeat that attains both.

B.1 Asynchronous SAC

Algorithm 1 Asynchronous SAC with Parallel Environments [Liu et al., 2025]

Require: Config d_{arg} , buffer size $|\mathcal{B}|$, LRs, γ , τ

- 1: Init shared queues `actor_queue`, `sample_queue`; stop signal `stop_event`
- 2: **Actor Process (CPU):** init $N_{\text{env}} = 13$ envs, local π_{θ_a}
- 3: **while** not `stop_event` **do**
- 4: Receive θ from `actor_queue` if available
- 5: Act: random if $t < T_{\text{start}}$, else π_{θ_a}
- 6: Step envs; push (s, a, r, s', d) to `sample_queue`
- 7: **end while**
- 8: **Learner Process (GPU):** init π_{θ} , $Q_{\phi_{1,2}}$, targets, $\mathcal{D}$
- 9: **while** total samples $< T$ **do**
- 10: Drain `sample_queue` into $\mathcal{D}$
- 11: **for** N_{loops} gradient steps **do**
- 12: Sample minibatch; compute SAC target y ; update Q , π , α ; soft-update targets
- 13: **end for**
- 14: Push updated θ to `actor_queue`
- 15: **end while**

451 B.2 Neural Architecture

452 **Scalar modes** utilize a Multi-Layer Perceptron (MLP):

$$\text{FC}(256) \xrightarrow{\text{Mish}} \text{FC}(256)$$

453 **Image modes** utilize an IMPALA-style CNN [Espeholt et al., 2018]. The feature extraction and
454 projection pipeline is defined as:

$$\begin{aligned} &\text{Conv}(n_c \rightarrow 32, k=8 \times 8, s=4) \xrightarrow{\text{Mish}} \\ &\text{Conv}(32 \rightarrow 64, k=4 \times 4, s=2) \xrightarrow{\text{Mish}} \\ &\text{Conv}(64 \rightarrow 64, k=3 \times 3, s=1) \xrightarrow{\text{Mish}} \text{Flatten} \rightarrow \text{FC}(256) \rightarrow \text{FC}(256) \end{aligned}$$

455 B.3 Hyperparameters

Table 2: SAC hyperparameters (TME v2 experiments).

Parameter	Symbol	Value
Agent decisions	T	1×10^5
Action repeat	–	6
Effective simulator steps	$6T$	6×10^5
Replay buffer size	$ \mathcal{B} $	2×10^5
Batch size	B	832
Warm-up steps	T_{start}	5,000
Policy / critic LR	η_π, η_Q	3×10^{-4}
Discount factor	γ	0.99
Target smoothing	τ	0.005
Reward weights	$w_{\text{cell}}, w_{\text{dose}}$	0.3, 2.0
Entropy autotuning	–	enabled
Parallel envs	N_{env}	13
Test frequency	–	every 4 episodes

456 B.4 Choice of action repeat

457 We select the action-repeat value with a dedicated sensitivity sweep, run independently of the main
458 observation-space experiments. We sweep `action_repeat` $\in \{1, \dots, 8\}$ and, for each value, measure
459 three end-of-run quantities aggregated over repeated rollouts (with 95% bootstrap CIs): tumor
460 reduction (control efficacy), an action-smoothness cost $\sum_t \|a_t - a_{t-1}\|^2$ (policy jitter; lower is
461 better), and a composite objective combining tumor reduction, dose, and smoothness (Figure 9).

462 Two effects drive the choice. (i) *Smoothness*: the action-smoothness cost falls steeply and mono-
463 tonically with action repeat from ≈ 110 at AR= 1 to ≈ 17 at AR= 6 and then flattens, so most of
464 the stabilisation benefit is already captured by AR= 6 and larger values add little. This matches
465 the physical timescale of the task: with a 15-minute decision interval on a slow immune process,
466 re-deciding every simulator step injects high-frequency drug-field jitter that does not change the
467 underlying dynamics. (ii) *Composite objective*: the objective is worst at AR= 1 (−0.88) and rises
468 to a plateau over AR= 5–8, with AR= 6 at −0.07; within that plateau the values are within each
469 other’s confidence intervals, so AR= 6 is chosen as the smallest repeat that reaches the objective
470 plateau while sitting at the knee of the smoothness curve. This makes action repeat a deliberate,
471 evidence-backed choice rather than an incidental hyperparameter.

472 C Full observation-mode comparison

473 Table 3 reports all nine observation modes evaluated, including the m1m2 scalar variants omitted from
474 the curated main-text Table 1. Train and test returns are end-of-training EWMA-50 means with the
475 standard deviation across seeds; train = rectangle, test = network-field.




Figure 9: Action-repeat sensitivity sweep (independent of the observation-space experiments; error bars are 95% bootstrap CIs). **Left:** tumor reduction (more negative = more cells removed). **Centre:** action-smoothness cost drops sharply and plateaus by $AR \approx 6$. **Right:** composite objective (higher is better) rises to a plateau over $AR = 5-8$; $AR = 6$ (dashed) is the smallest repeat on that plateau and at the knee of the smoothness curve.

Table 3: All observation modes (seeds averaged; train = rectangle, test = network-field). Modes are ordered by held-out test return.

ID	Mode	n	Train $_{\mu}$	Train $_{\sigma}$	Test $_{\mu}$	Test $_{\sigma}$
I2m	img_mc_cells_substrates_m1m2	5	+40.8	8.5	+32.5	9.4
I1	img_mc_cells	6	+30.5	8.7	+32.4	10.5
I2	img_mc_cells_substrates	7	+42.5	6.6	+30.4	7.4
I1m	img_mc_cells_m1m2	5	+35.0	13.1	+26.0	12.4
S5m	spatial_scalars_cells_spatial_no_scalars_substrates_m1m2	5	+53.6	34.2	+9.5	9.6
S3sm	spatial_scalars_cells_substrates_m1m2	5	+38.3	27.8	+2.6	12.1
S3m	spatial_scalars_cells_m1m2	6	+45.7	26.8	+2.6	9.7
S3s	spatial_scalars_cells_substrates	5	+42.9	18.8	-0.0	9.1
RAND	random policy (baseline)	5	-18.4	4.0	-25.3	3.0
POMDP	scalars_macrophages	5	-20.4	8.3	-31.8	2.3

D Bootstrap confidence intervals

Because each mode is trained with only $n = 5-7$ seeds, the across-seed standard deviations reported in Tables 1 and 3 are themselves noisy estimates of dispersion. To quantify uncertainty on the *mean* return more robustly, we bootstrap over seeds: for each mode we resample the seeds with replacement $B = 10,000$ times, recompute the mean end-of-training return of each resample, and report the 2.5th–97.5th percentiles of that bootstrap distribution as a 95% confidence interval on the mean (Table 4). The same procedure applied per step yields the shaded bands in Figures 10 and 11.

The bootstrap intervals sharpen the main-text claims. On the held-out network-field, the three image modes have test-return CIs that lie entirely above +24, whereas every spatial-scalar mode has a test-return CI that *contains zero* (e.g. S3s $[-6.2, +9.1]$), so their generalization is not statistically distinguishable from a non-learning policy. The image and spatial-scalar CIs do not overlap, confirming that the generalization gap is significant and not an artefact of seed selection. The scalars_macrophages POMDP baseline has a tight, strongly negative CI $[-33.4, -29.4]$, reflecting consistent failure rather than robustness and it does not overlap the random policy’s test CI $[-28.2, -22.8]$ from the favourable side, so the macrophage-only observation performs *no better than, and plausibly worse than, random dosing*.

E Example Episode Comparisons

Figure 12 shows a single matched test episode comparing an image mode (img_mc_cells_substrates), a spatial-scalar mode (spatial_scalars_cells_substrates), and the POMDP baseline (scalars_macrophages) step by step. To guarantee an exact comparison,

Table 4: Mean end-of-training return with a 95% bootstrap confidence interval on the mean ($B = 10,000$ seed resamples; train = rectangle, test = network-field). Ordered by test mean.

ID	Mode	n	Train $_{\mu}$ (95% CI)	Test $_{\mu}$ (95% CI)
I2m	img_mc_cells_substrates_m1m2	5	+40.8 [+33.0, +47.3]	+32.5 [+24.4, +40.6]
I1	img_mc_cells	6	+30.5 [+23.1, +36.7]	+32.4 [+24.2, +40.6]
I2	img_mc_cells_substrates	7	+42.5 [+37.4, +47.2]	+30.4 [+24.5, +36.3]
I1m	img_mc_cells_m1m2	5	+35.0 [+22.7, +45.6]	+26.0 [+13.5, +38.5]
S5m	spatial_scalars_cells_spatial_no_scalars_substrates_m1m2	5	+53.6 [+25.0, +85.1]	+9.5 [+0.4, +17.6]
S3sm	spatial_scalars_cells_substrates_m1m2	5	+38.3 [+20.1, +65.9]	+2.6 [-6.2, +11.4]
S3m	spatial_scalars_cells_m1m2	6	+45.7 [+24.9, +67.3]	+2.6 [-5.8, +10.9]
S3s	spatial_scalars_cells_substrates	5	+42.9 [+24.7, +58.8]	-0.0 [-6.2, +9.9]
RAND	random policy (baseline)	5	-18.4 [-21.9, -14.9]	-25.3 [-28.2, -22.4]
POMDP	scalars_macrophages	5	-20.4 [-28.0, -13.8]	-31.8 [-33.4, -30.2]




Figure 10: Training return on rectangle (EWMA-50, mean with 95% bootstrap CI band over seeds) vs. cumulative simulator steps. Curated main modes only.

all three agents are rolled out (deterministically, using the policy mean action) from the *identical* held-out network-field initial condition the same 124-tumor-cell layout replayed from file so the only difference between the three trajectories is the observation mode.

The episode cleanly separates the three regimes. `img_mc_cells_substrates` drives the tumor to near-eradication (3 cells remaining, +223 cumulative return); it spends the most drug (4.5 total dose) but places it precisely enough visibly dosing in sustained, targeted bursts (Figure 12, bottom) to convert dose into effective macrophage repolarisation and tumor clearance. `spatial_scalars_cells_substrates` achieves partial control (70 cells, +30 return, 2.8 dose): it suppresses growth early but plateaus and cannot finish, lacking the spatial substrate feedback needed to target the final push. `scalars_macrophages` (POMDP) barely acts (0.9 total dose) and the tumor *grows* past its starting size (130 cells, -6.4 return): seeing only macrophage-polarisation counts, it cannot localise where to inject, so it under-doses and fails consistent with its aggregate performance at or below the random baseline.




Figure 11: Held-out test return on network-field (EWMA-50, mean with 95% bootstrap CI band over seeds) vs. cumulative simulator steps. The image modes' bands separate from the spatial-scalar summary (S3s), whose band hugs zero, and from the POMDP baseline.

Matched network-field episode (identical IC, 124 tumor cells)




Figure 12: **Matched network-field episode (identical initial condition, 124 tumor cells)**. Cumulative reward (top), tumor cell count (middle), and the dose action (bottom, rolling mean over 6 steps matching the action-repeat horizon used in training; the raw per-step signal is noisy and unreadable at full resolution) for `img_mc_cells_substrates` (red), `spatial_scalars_cells_substrates` (blue), and the `scalars_macrophages` POMDP baseline (grey), all replayed from the same IC. `img_mc_cells_substrates` reaches near-eradication (3 cells, +223 return, 4.5 dose); `spatial_scalars_cells_substrates` achieves partial control (70 cells, +30 return, 2.8 dose); the POMDP baseline under-doses and the tumor grows (130 cells, -6.4 return, 0.9 dose). Only the image mode sustains a high, decisive dose action long enough to drive the tumor to near-eradication, whereas `spatial_scalars_cells_substrates` oscillates and eventually lets the tumor regrow. Note that the physical dose delivered to the simulation each step is the concentration action (bottom) multiplied by the voxel volume swept by the independently-controlled injection radius action, so the dose action alone does not linearly predict the physical dose spent.